




BASELINE MODEL

| Experiment # | Epochs | Learning Rate | Batch Size | Optimizer | Test Accuracy | Average Loss |
|--------------|--------|---------------|------------|-----------|---------------|--------------|
| 1 | 20 | 0.1 | 64 | SGD | 76.6% | 0.51 |
| 2 | 20 | 0.01 | 64 | SGD | 83.6% | 0.33 |
| 3 | 20 | 0.001 | 64 | SGD | 83.6% | 0.42 |
| 4 | 20 | 0.0001 | 64 | SGD | 80.5% | 0.82 |
| 5 | 20 | 0.001 | 16 | SGD | 83.9% | 0.47 |
| 6 | 20 | 0.001 | 32 | SGD | 83.9% | 0.53 |
| 7 | 20 | 0.001 | 64 | SGD | 83.6% | 0.46 |
| 8 | 20 | 0.001 | 128 | SGD | 83.2% | 0.50 |
| 9 | 20 | 0.001 | 512 | SGD | 81.0% | 0.92 |
| 10 | 20 | 0.001 | 1024 | SGD | 78.9% | 1.34 |
| 11 | 1 | 0.001 | 64 | SGD | 78.1% | 0.78 |
| 12 | 10 | 0.001 | 64 | SGD | 83.1% | 0.53 |
| 13 | 20 | 0.001 | 64 | SGD | 83.6% | 0.46 |
| 14 | 50 | 0.001 | 64 | SGD | 83.9% | 0.47 |
| 15 | 100 | 0.001 | 64 | SGD | 84.0% | 0.43 |
| 16 | 20 | 0.001 | 64 | Adam | 82.9% | 0.34 |
| 17 | 20 | 0.001 | 64 | SGD | 83.6% | 0.46 |






FINDINGS

- * The learning rate significantly impacts model performance. Too high or too low rates hurt performance.
- * Smaller batches are better. Batch sizes of 16-64 perform well (83.6-83.9% accuracy), while larger batches (512-1024) drop accuracy and increase loss significantly.
- * More epochs help, but only up to a point. Accuracy improves from 78.1% at 1 epoch to 83.6% at 20 epochs. After that, the gains are minimized.
- * Using Adam optimizer instead of SGD results in slightly lower accuracy (82.9%) but a lower average loss (0.34). This means that Adam may converge faster or more smoothly but does not necessarily mean it's better than SGD in this case.

IMPROVED MODEL

| Experiment # | Epochs | Learning Rate | Batch Size | Optimizer | Test Accuracy | Average Loss |
|--------------|--------|---------------|------------|-----------|---------------|--------------|
| 1 | 20 | 0.1 | 64 | Adam | 0.4% | 6.04 |
| 2 | 20 | 0.01 | 64 | Adam | 72.5% | 0.94 |
| 3 | 20 | 0.001 | 64 | Adam | 91.1% | 0.11 |
| 4 | 20 | 0.0001 | 64 | Adam | 91.3% | 0.31 |
| 5 | 20 | 0.001 | 16 | Adam | 89.9% | 0.18 |
| 6 | 20 | 0.001 | 32 | Adam | 90.6% | 0.14 |
| 7 | 20 | 0.001 | 64 | Adam | 91.3% | 0.09 |
| 8 | 20 | 0.001 | 128 | Adam | 91.3% | 0.07 |
| 9 | 20 | 0.001 | 512 | Adam | 90.8% | 0.23 |
| 10 | 20 | 0.001 | 1024 | Adam | 91.2% | 0.29 |
| 11 | 1 | 0.001 | 64 | Adam | 84.1% | 0.35 |
| 12 | 10 | 0.001 | 64 | Adam | 90.5% | 0.14 |
| 13 | 20 | 0.001 | 64 | Adam | 92.0% | 0.10 |
| 14 | 50 | 0.001 | 64 | Adam | 91.8% | 0.08 |
| 15 | 100 | 0.001 | 64 | Adam | 92.0% | 0.06 |
| 16 | 20 | 0.001 | 64 | SGD | 91.6% | 0.29 |
| 17 | 20 | 0.001 | 64 | Adam | 92.0% | 0.10 |



FINDINGS

- * A learning rate of 0.001 works best, achieving 91.1-92.0% accuracy. Higher rates (e.g., 0.1) fail completely (0.4% accuracy), while lower rates (0.0001) still perform well but with slightly higher loss.
- * Batch size has minimal effect on accuracy, with all sizes (16-1024) achieving 89.9-91.3% accuracy.
- * Increasing epochs from 1 to 20 significantly improves accuracy (84.1% to 92.0%) but beyond 20 epochs, gains are minimal. Accuracy plateaus around 92.0%.
- * Adam achieves higher accuracy (92.0%) and lower loss (0.10) compared to SGD (91.6% accuracy, 0.29 loss).